

# Lecture 3 – Processes and the Operating System

---

# Outline

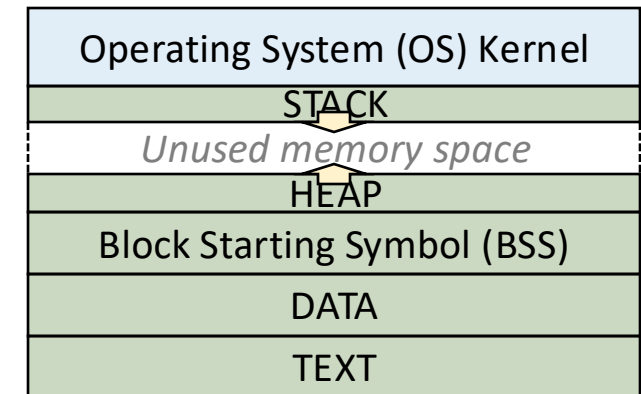
---

- Recap of computer architecture
- What is the OS?
  - Kernel, process
- CPU Virtualisation
- System Calls (remind what is a CPU exception)
- How do you execute a process over an OS?
- Scheduling

# Computer Architecture Recap

- Instruction-set architectures is the lowest **interface between computer hardware and software**
  - It defines *how software tells the hardware what to do*
  - Defined by the CPU architecture
  - Examples: ARM Assembly
- Most of our program assumed exclusive access to CPU resources and no OS was used.
  - What happens if two assembly processes want to access the CPU?


31	0
R0	
R1	
R2	
R3	
R4	
R5	
R6	
R7	
R8	
R9	
R10	
R11	
R12	
SP (R13)	
LR (R14)	
PC (R15)	



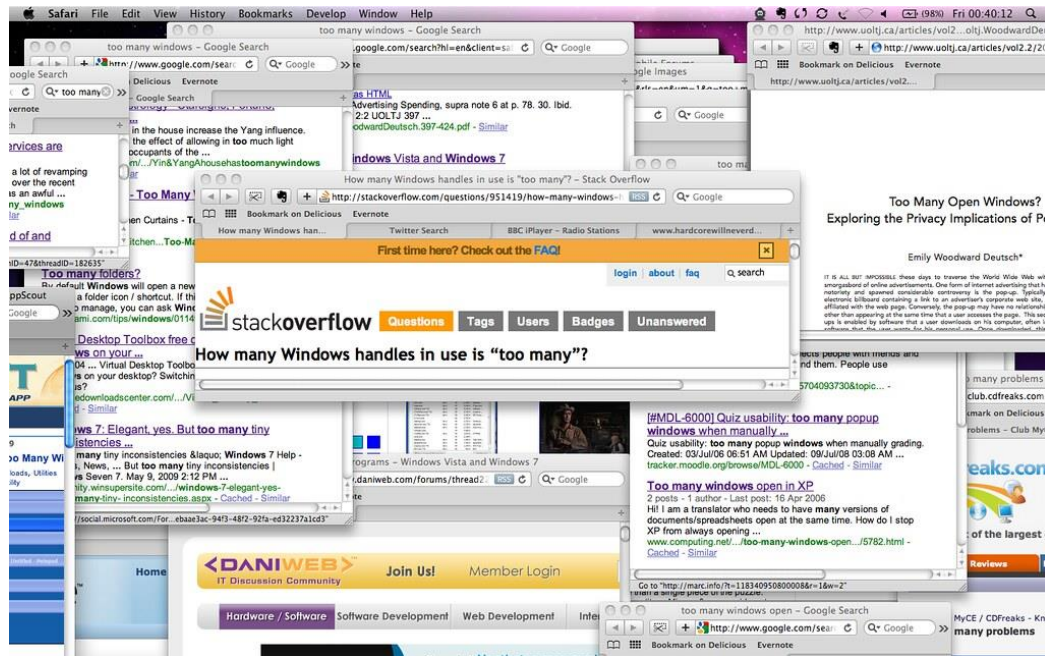
Instructions	
⋮	⋮
0040000C	E 3 A 0 1 0 6 4
00400008	E 3 A 0 2 0 4 5
00400004	E 1 5 1 0 0 0 2
00400000	2 5 8 1 3 0 2 4 ← PC
⋮	⋮

# Single-process computers

Single program loaded in memory, executed to completion



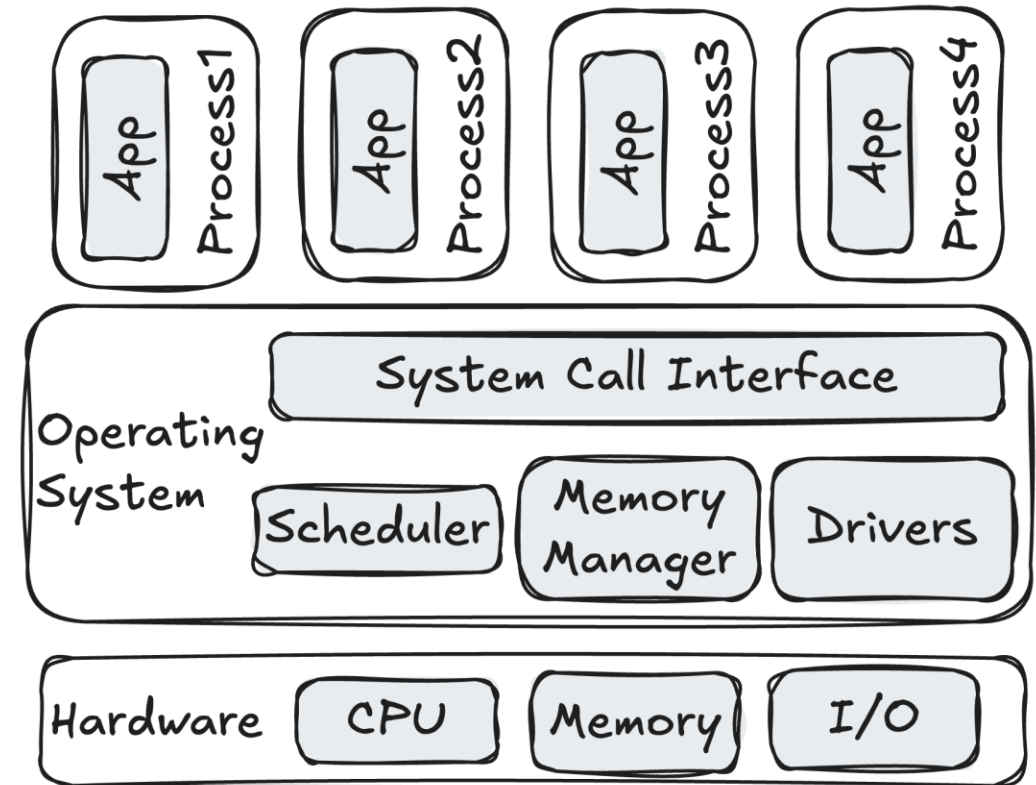
# Modern computers



- Lots of processes running in parallel...
- Lots of IO devices shared...
- Programs are responsive...
- How do modern computers achieve parallelism?

# What is an operating system?

- An **Operating System** is a software layer that sits between hardware and user applications.
  - Manage resources
  - Safe and convenient programming abstractions
  - Services: scheduling, device drivers
- CPU virtualization: control which process uses the CPU



# Virtualization

---

- A combination of indirection and multiplexing
- Create a **virtual version of a resource** (hardware, OS, network, or storage) so that multiple independent systems or applications can share it as if each has its own dedicated resource
- Examples:
  - virtual memory
  - virtual modem
  - Cloud computing
  - Santa claus
- Can cleanly and dynamically reconfigure the system

# OS Virtualisation

---


- **Goal:** give each process the illusion of exclusive resource access
- **Constraints:** The resources are fixed and shared among all processes, ensuring isolation and security
- **Time-sharing:** exclusive use, one at a time
- **Space-sharing:** everyone gets a small chunk all the time
- Different strategies for resources
  - CPU: time sharing, alternate between tasks
  - Memory: space sharing (more later)
  - Disk: space sharing (more later)



# What is a process?

---


- A **program** is a bundle of binary instructions and static data, stored on the disk.
- A **process** is an OS abstraction that represents a program loaded into memory and executed. It includes:
  - Memory resources
  - Execution state (registers)
  - Resources (e.g., open files)
  - Process control block (state maintained by the OS)



# What is a process?

---

- A **program** is a bundle of binary instructions and static data, stored on the disk.
- A **process** is a running instance of a program. It includes:
  - Memory resources
  - Execution state (registers)
  - Resources (e.g., open files)
  - Process control block (state maintained by the OS)
- A process can have one or more **threads**.



# Process API

---

- UNIX systems offer an extended API system to manage processes:
- *fork()*: duplicate the current process (copy memory and register state)
  - The two processes will have different PID
  - The original process is now the *parent process*, and the new process is a *child process*.
- *exec()* / *execvp()*: replaces the current process's memory image with a new program
- *wait()*: parent process waits for one of its child processes to finish
  - If wait is not called, the child process enters zombie state (terminated, but the state not thoroughly cleaned up)
- Lots of examples during the lab activities this week...

# Process Creation

---

- OS allocates internal data structures
- OS allocates an address space
  - Loads code, data from disk
  - Creates runtime stack, heap
- OS opens basic files (STDIN, STDOUT, STDERR)
- OS initializes CPU registers

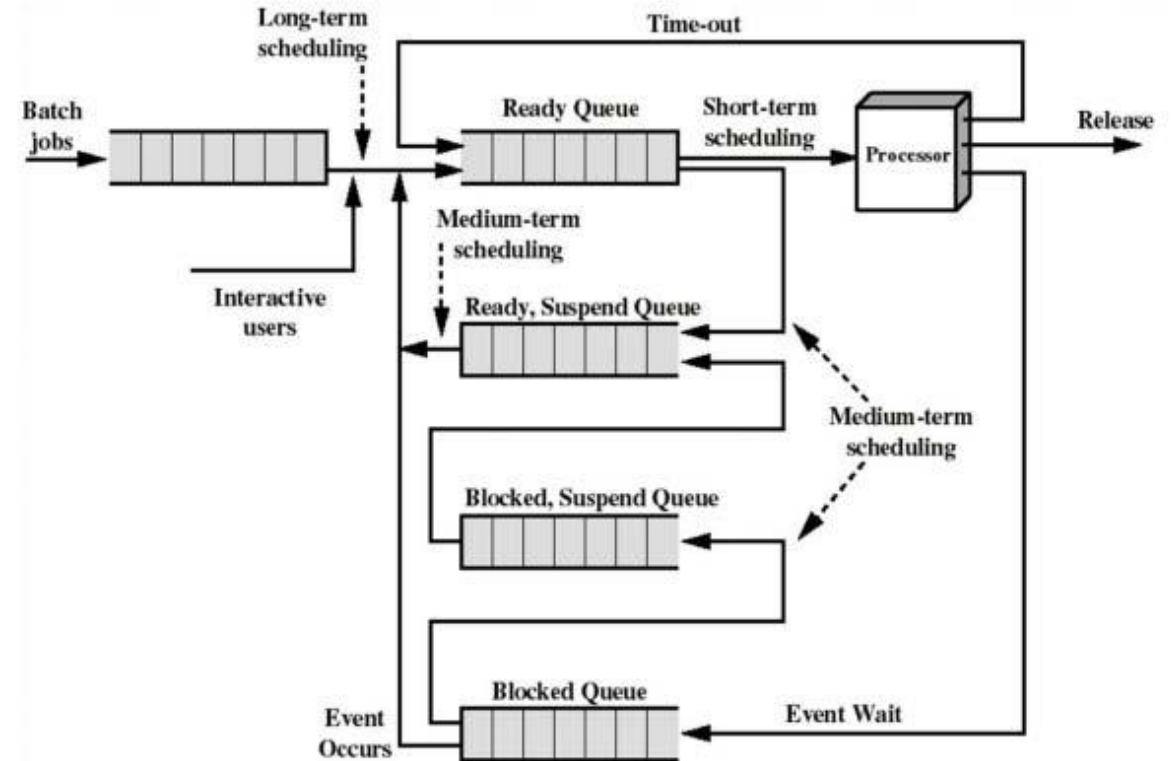
# CPU virtualisation

---

- When the user executes a program, the OS creates a process.
- OS time-shares CPU across multiple processes.
  - OS scheduler picks one of the executable processes to run.
  - Scheduler must keep a list of processes
  - Scheduler must keep metadata for policy
- **How do you execute code in a secure way, without incurring performance penalties?**

# OS Scheduling

- An **OS scheduler** is the part of the operating system that decides **which process gets to use the CPU (or other resources) at any given time.**
  - It ensures efficient, fair, and responsive execution of multiple processes.
- Extended research explores how to optimize different performance metrics:
  - Throughput, waiting time, response time, fairness...
- More details in SCC.233...



# Direct Execution (no limits)

OS	Program
<i>Create entry for process list</i> <i>Allocate memory for program</i> <i>Load program into memory</i> <i>Set up stack with argc/argv</i> <i>Clear registers</i> <i>Execute call main()</i>	
	<i>Run main()</i> <i>Execute return from main</i>
<i>Free memory of process</i> <i>Remove from process list</i>	


*Problem 1 (security): if we just run a program, how can the OS make sure the program doesn't do anything that we don't want it to do, while still running it efficiently?*

*Problem 2 (sharing): when we are running a process, how does the operating system stop it from running and switch to another process, thus implementing the time sharing we require to virtualize the CPU?*

# Problem 1: Security

---

- A process needs to access the disk.
  - How do you check user permissions?
- Solution: protection rings enforced by the CPU
  - User mode: the code cannot access devices and kernel data (processes)
  - Kernel model: the code has unlimited access to resources (kernel)
- How can a process request a service from the kernel (e.g., access a file?)
  - Special type of exception, called **system call**.





# Exceptions

---

- The programmer defines program execution.
  - Nonetheless, events occur external to the program logic.
- An **exception** is used to model an unexpected event during program execution.
  - It looks like a function call that branches to a new address.
  - Allows a program/OS to implement code that manages external events.
- Exceptions may be caused by hardware or software.
  - User presses a key on a keyboard (hardware exception, input/output (I/O) device);
  - Encounter undefined instructions.
  - Trap instruction (SVC - ARM, INT – x86); a program invokes a service OS, running at a higher privilege level.

# System call

---

- A **system call** allows programs to request a service from the kernel.
  - It's the **interface** between user programs (in user space) and the OS (in kernel space).
  - Linux offers a couple of 100's system calls: open, write, read, exit ...
- The OS at boot registers the address of a system call handler function with the CPU
- A process executes the trap instruction
- The process context is stored in memory, the state of the OS is loaded and the CPU executes the call handler function (see previous slide)

# x86 System Calls

Show 10 entries										Search: <input type="text"/>	
#	Name	Registers								Definition	
		eax	ebx	ecx	edx	esi	edi	ebp			
0	restart_syscall	0x00	-	-	-	-	-	-	kernel/signal.c:2501		
1	exit	0x01	int error_code	-	-	-	-	-	kernel/exit.c:1095		
2	fork	0x02	-	-	-	-	-	-	arch/x86/kernel/process.c:271		
3	read	0x03	unsigned int fd	char *buf	size_t count	-	-	-	fs/read_write.c:460		
4	write	0x04	unsigned int fd	const char *buf	size_t count	-	-	-	fs/read_write.c:477		
5	open	0x05	const char *filename	int flags	umode_t mode	-	-	-	fs/open.c:1046		
6	close	0x06	unsigned int fd	-	-	-	-	-	fs/open.c:1117		
7	waitpid	0x07	pid_t pid	int *stat_addr	int options	-	-	-	kernel/exit.c:1879		
8	creat	0x08	const char *pathname	umode_t mode	-	-	-	-	fs/open.c:1079		
9	link	0x09	const char *oldname	const char *newname	-	-	-	-	fs/namei.c:3152		
10	unlink	0x0a	const char *pathname	-	-	-	-	-	fs/namei.c:2979		
11	execve	0x0b	const char *name	const char *const *argv	const char *const *envp	-	-	-	arch/x86/kernel/process.c:342		
12	chdir	0x0c	const char *filename	-	-	-	-	-	fs/open.c:375		
13	time	0x0d	time_t *tloc	-	-	-	-	-	kernel/time.c:62		
14	mknod	0x0e	const char *filename	umode_t mode	unsigned dev	-	-	-	fs/namei.c:2693		
15	chmod	0x0f	const char *filename	umode_t mode	-	-	-	-	fs/open.c:499		
16	lchown	0x10	const char *filename	uid_t user	gid_t group	-	-	-	fs/open.c:586		

<https://syscalls.w3challs.com/?arch=x86>

## Problem 2: CPU sharing

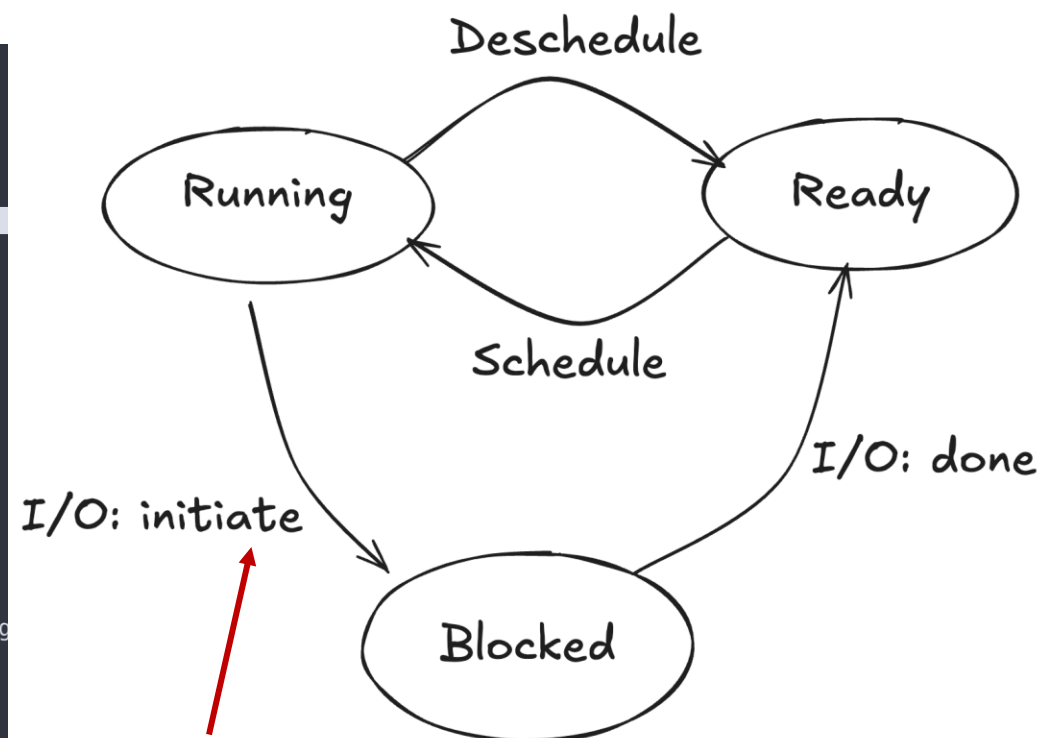
---

- If a process runs (i.e., uses the CPU), the OS cannot context switch to share the CPU.
  - *How can the operating system regain control of the CPU so that it can switch between processes?*
- Cooperative scheduling: Long-running processes periodically give up the CPU for the OS to run other tasks.
- Non-cooperative schedule: Schedule the OS over fixed intervals (scheduling quantas) using a timer interrupt.

# Process states

```
top - 12:22:24 up 85 days, 21:26, 4 users, load average: 0.36, 0.20, 0.13
Tasks: 344 total, 1 running, 343 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.2 us, 0.1 sy, 0.0 ni, 99.6 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 31779.4 total, 1399.9 free, 2895.6 used, 27483.9 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 27698.5 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1290	root	20	0	3118700	59168	24968	S	6.0	0.2	5816:04	kubelet
20879	root	10	-10	2005512	235888	12528	S	1.3	0.7	10:01.74	ovs-vswitchd
21040	root	20	0	6687464	902008	28176	S	1.3	2.8	36:33.54	qemu
20161	root	20	0	4191320	126976	63104	S	1.0	0.4	49:41.41	dockerd
5688	root	20	0	1257972	43348	24204	S	0.3	0.1	0:00.53	tailscaled
6282	tudor	20	0	11088	4632	3752	R	0.3	0.0	0:00.24	top
20692	root	20	0	1597720	147884	106016	S	0.3	0.5	8:07.01	grafana
20721	root	20	0	2710932	80904	49676	S	0.3	0.2	151:22.39	promtail
1	root	20	0	381372	14112	8552	S	0.0	0.0	178:58.70	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:16.38	kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_gp
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_par_gp
5	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	slub_flushwq
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	netns
8	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:0H-events_highpri
10	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_wq
11	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_rude
12	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_tasks_trace
13	root	20	0	0	0	0	S	0.0	0.0	6:10.41	ksoftirqd/0



Processes can block because they expect IO to complete from a computer device.

# Limited direct execution (Boot)

---

OS @ boot	Hardware	Program
initialize trap table	remember address of... syscall handler timer handler	
Start interrupt timer	Start timer Interrupt CPU in X msec	

# Limited direct execution (Start process)

OS@ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg/PC return-from-trap	restore regs (from kernel stack) move to user mode jump to main	Run main() ...

# Limited Direct Execution (System Call)

OS@ run (kernel mode)	Hardware	Program (user mode)
Handle trap Do work of syscall <b>return-from-trap</b>	save regs (to kernel stack) move to kernel mode, jump to trap handler  restore regs (from kernel stack) move to user mode jump to PC after trap	Call system call trap into OS      ... return from main trap (via exit())
Free memory of process Remove from process list		



# Limited Direct Execution (Timer interrupt)

OS@ run (kernel mode)	Hardware	Program (user mode)
<p>Handle the trap Call <code>switch()</code> routine   save <code>regs(A)</code> to <code>proc-struct(A)</code>   restore <code>regs(B)</code> from <code>proc-struct(B)</code>   switch to <code>k-stack(B)</code> <b>return-from-trap (into B)</b></p>	<p><b>timer interrupt</b> save <code>regs(A)</code> to <code>k-stack(A)</code> move to kernel mode, jump to timer handler</p> <p><i>restore <code>regs(B)</code> from <code>k-stack(B)</code> move to user mode, jump to B's PC</i></p>	<p>Process A</p> <p>Process B</p>

# Conclusion

---

- An **Operating System** is a software layer that sits between hardware and user applications.
  - Manage resources
  - Safe and convenient programming abstractions
  - Services: scheduling, device drivers
- CPU virtualization allows multiple processes to share a physical CPU.
- Limited direct code execution allows the OS scheduler to control process scheduling.
  - Processes use system calls to access virtualized resources.
  - Timer interrupts enable the OS to schedule processes.
- Next: Memory Virtualisation

# Further reading

---

- Chapter 4: Processes
- Chapter 5: Process API
- Chapter 6: Direct Execution

