

Lecture 2 – Introduction to System Design

Outline

What is system design?


- Optimisation
- Bottleneck
- Constraint
- Resource
- metric

System design patterns:

- Multiplexing
- Pipelining
- Batching
- Locality/Caching
- Distributed/Hierarchy
- Abstraction/Binding and Indirection
- Virtualisation / Randomisation
- Data and control separation


What is system design?

- A **system** is a collection of **components** (such as hardware and software) that work together to achieve a **specific goal**.
 - A computer consists of CPU, memory, devices
 - A network consists of end-hosts, links, switches, routers
- **System design** is the science (and a bit of art) of putting together these resources into a harmonious whole
 - *Extract the most from what you have (optimisation)*
- *In our module, we will explore how system design techniques can be applied in OS, networking and the cloud.*



System Optimization

- In any system, some resources are more freely available than others
 - High-end PC connected to the Internet via a poor and noisy WiFi link
 - *The constrained* resource is link bandwidth
 - PC, CPU, and memory are *unconstrained*
- **Optimization: Maximize a set of performance metrics (e.g., time to load webpage) given a set of resource constraints (e.g. bandwidth)**
- Explicitly identifying *metrics*, *resources*, and *constraints* helps in designing efficient systems



Common Resources (1)

- **Time:** can be translated into multiple aspects
 - deadline for task completion
 - time to market
 - mean time between failures
- **Space:** Shows up as
 - Number of PCI slots
 - limit to available memory (kilobytes)
 - bandwidth (kilobits)
 - 1 kilobit/s = 1000 bits/sec, but 1 kilobyte/s = 1024 bits/sec!

Common Resources (2)

- **Computation:** Amount of processing that can be done in unit time
 - Can increase computing power by
 - using more processors
 - waiting for a while!
- **Cost:** cost to implement system
 - what components can be used
 - what price users are willing to pay for a service
 - the number of engineers available to complete a task
- **Labor:** Human effort required to design and build a system
 - Constrains what can be done, and how fast

Constraints: Scaling

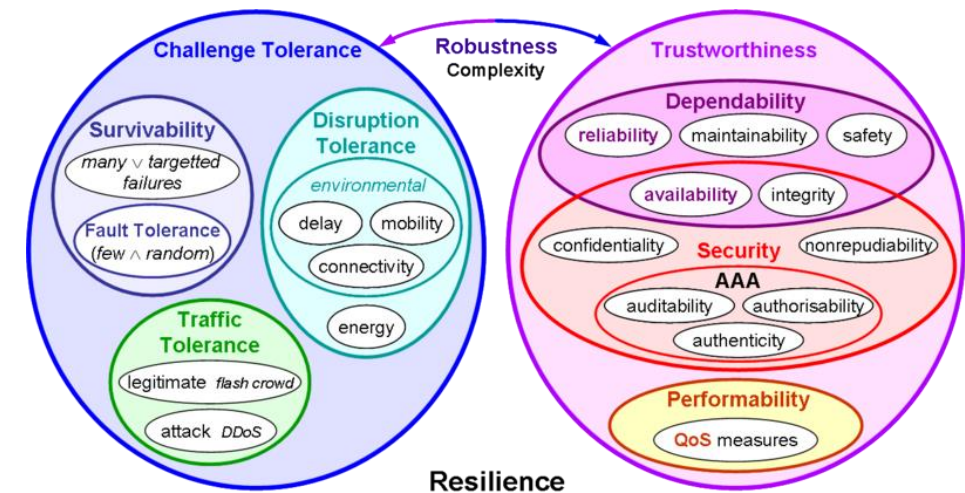
- **Scaling** is the ability of a system to handle an increasing amount of work, users, or data by adding resources (such as CPU, memory, bandwidth, or servers) without a proportional loss of performance.
 - A design constraint, rather than a resource constraint
- Can use any centralized elements in the design
 - forces the use of complicated distributed algorithms
- Hard to measure
 - but necessary for success

Constraints: Social

- **Standards**
 - force design to conform to requirements that may or may not make sense
 - underspecified standard can faulty and non-interoperable implementations
- **Market requirements**
 - products may need to be backwards compatible
 - may need to use a particular operating system
 - example
 - GUI-centric design

Constraints: Resilience

- **Resilience** is the ability of a system to **continue operating correctly, or recover quickly**, when it faces unexpected problems such as hardware failures, software bugs, high load, or cyberattacks
 - Includes quantitative (availability) and qualitative aspects (safety)
- Typically, resilience depends on risk analysis/threat model to understand constraints.




Metrics

Category	Goal / What It Measures	Typical Metrics
Performance / Efficiency	Speed or responsiveness of the system	Latency (response time) Throughput (ops/sec) Utilization (CPU, link)
Scalability	How performance changes with workload or resources	Elasticity (load adaptation) Parallel efficiency
Reliability / Availability / Resilience / Security	Ability to function despite failures	MTBF (Mean Time Between Failures) Availability = $\text{Uptime} / (\text{Uptime} + \text{Downtime})$ Isolation strength
Cost / Resource Utilization	Efficiency in using limited resources	CapEx / OpEx - Memory /Network bandwidth usage
Maintainability / Evolvability	Ease of modification, debugging, or scaling	Code complexity (LOC, cyclomatic) Deployment time - Configuration overhead
Sustainability	Environmental, power and material impact	Energy per operation Carbon footprint (CO2 mass)

System Bottlenecks

- A **bottleneck**, is the most constrained element in a system with respect to our optimization metric (e.g. longest/slowest process)
- Performance metrics improves by removing system bottlenecks
 - This process can create new bottlenecks
- In a *balanced* system, all resources are simultaneously bottlenecked
 - this is optimal
 - but nearly impossible to achieve
 - in practice, bottlenecks move from one part of the system to another



System Design in Real Life

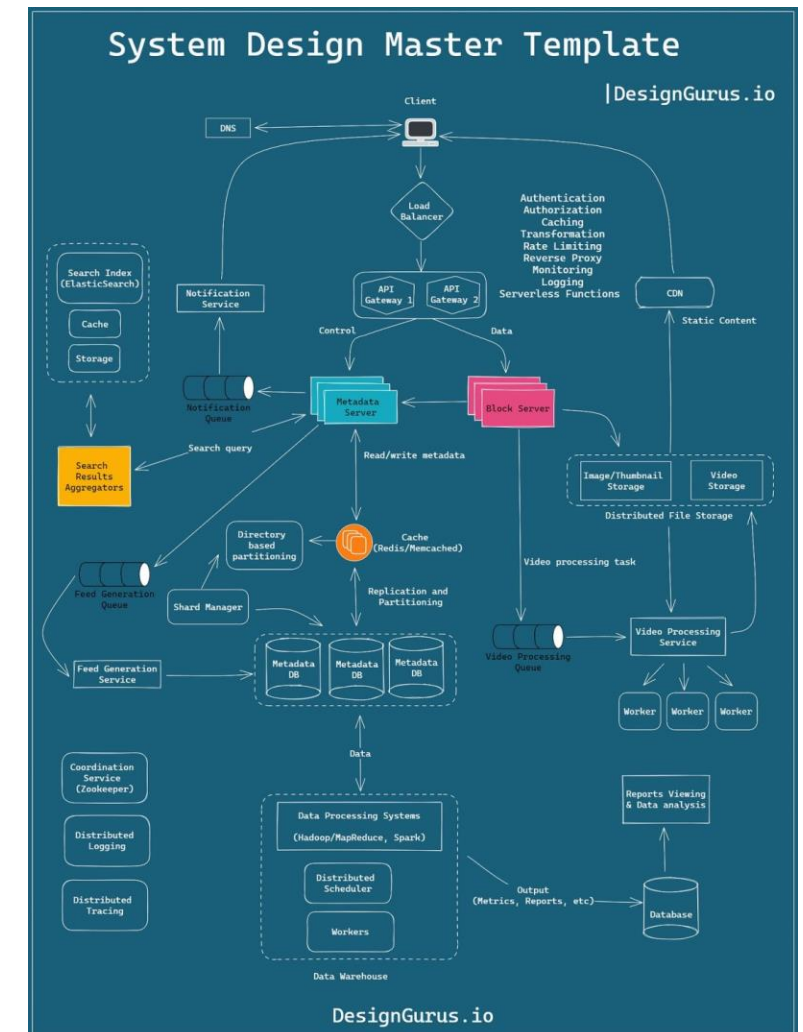
- Can't always quantify and control all system metrics
 - Metrics such as *scalability*, *modularity*, *extensibility*, and *elegance* are essential, but unquantifiable
- Rapid technological change can add or remove resource constraints
 - Multi-core CPUs can overcome the limitations of CPU clock speed, but need a rethink for application architectures
- Market conditions may dictate changes to the design halfway through the process
- Nevertheless, still possible to identify some principles



System Desing Patterns




- Use unconstrained resources to alleviate bottleneck
- Several **standard design patterns** allow us to trade off one resource for another
- *We will revisit these theoretical concepts and explain how they apply in different scenarios in OS and network design.*



Multiplexing

- Use the same resource to service multiple requests without interference
 - Another word for sharing
- Goal: serve a large number of users, without increasing resource requirements.
 - Excellent for coping with **cost** constraints
 - Users see an increased response time, and take up space when waiting, but the system costs less
 - economies of scale



Multiplexing examples


- Multiplexing dimensions:
 - Time: CPU and I/O multiplexing
 - Multiple processes share the CPU and read/write on the same disk
 - OS scheduler ensure processes do not interfere
 - Space: Memory multiplexing
 - Virtual memory gives each process its own address space
 - Frequency: Radio comms
 - Mobile networks use the same airwaves to transmit signal.
 - Complex algorithms use multiple frequencies to increase user number, without affecting performance and resilience.

Statistical Multiplexing

- Share resources dynamically based on demand.
- Suppose the resource has capacity C
- Shared by N identical tasks
- Each task requires capacity c
- If $Nc \leq C$, then the resource is underloaded
- If at most 10% of tasks are active, then $C \geq Nc/10$ is enough
 - We use the statistical knowledge of users to reduce system cost
 - This is *the statistical multiplexing gain*


Example of statistical multiplexing gain

- Consider a 100-room hotel
- How many external phone lines does it need?
 - Trade-off: each line costs money to install and rent
- What if a voice call is active only 40% of the time?
 - Rent only 40 lines, drop call if > 40 simultaneous calls
 - Negatively impact on user satisfaction, if peak load estimation is low
- Good statistics on resource use are essential to achieve gains
 - If statistics are incorrect or change over time, we're in trouble
 - Example: road system



Pipelining (1)

- Suppose you have a time constraint; i.e. complete a task in less time
 - Could you use more processors to do so?
- Yes, if you can break up the task into *independent* subtasks
 - Example: downloading page objects in a browser
 - optimal if all subtasks take the same time
- What if subtasks are dependent?
 - for instance, a subtask may not begin execution before another ends
 - such as in cooking
- Then, having more processors doesn't always help.

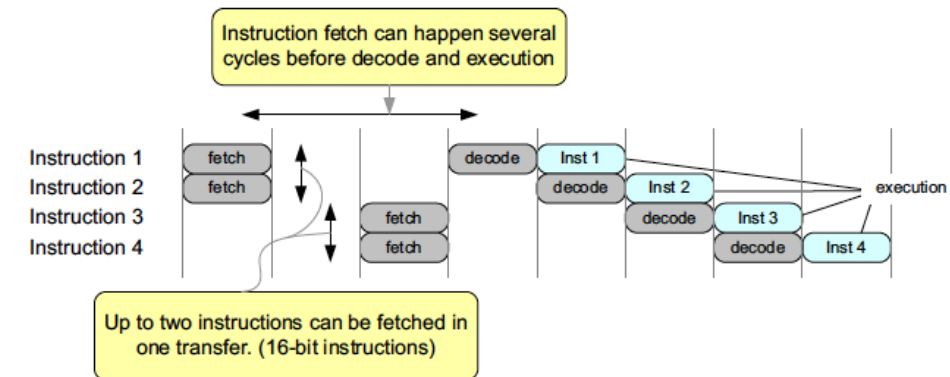


Pipelining (2)

- Special case of *serially dependent* subtasks
 - A subtask depends only on the previous one in the execution chain
- Can use a *pipeline*
 - think of an assembly line
- In pipelining, multiple **stages** overlap during execution.
 - Each processor works on a different stage of the process.
- What is the best decomposition?
 - If the sum of times taken by all stages = R
 - Slowest stage takes time S
 - Throughput = $1/S$
 - Response time = R
 - Degree of parallelism = R/S
 - Number of processors = N
 - Maximize parallelism when $R/S = N$, so that $S = R/N \Rightarrow$ equal stages
 - *balanced pipeline*



Batching

- Group tasks together and process them as a single unit, rather than handling them individually.
 - Amortise overheads (e.g.. loading time)
- Only works when *`overhead for N tasks` < `N time overhead for one task`* (i.e. *nonlinear*)
 - If loading two 16-bit instructions takes the same time to load 1 16-bit instruction, then you improved read access times.
- Also, the time taken to accumulate a batch shouldn't be too long
- We're trading off reduced overhead for a longer worst-case response time and increased throughput
 - Improve average response **time** of the system.
 - Improve CPU constraints; Complete more instructions per second.



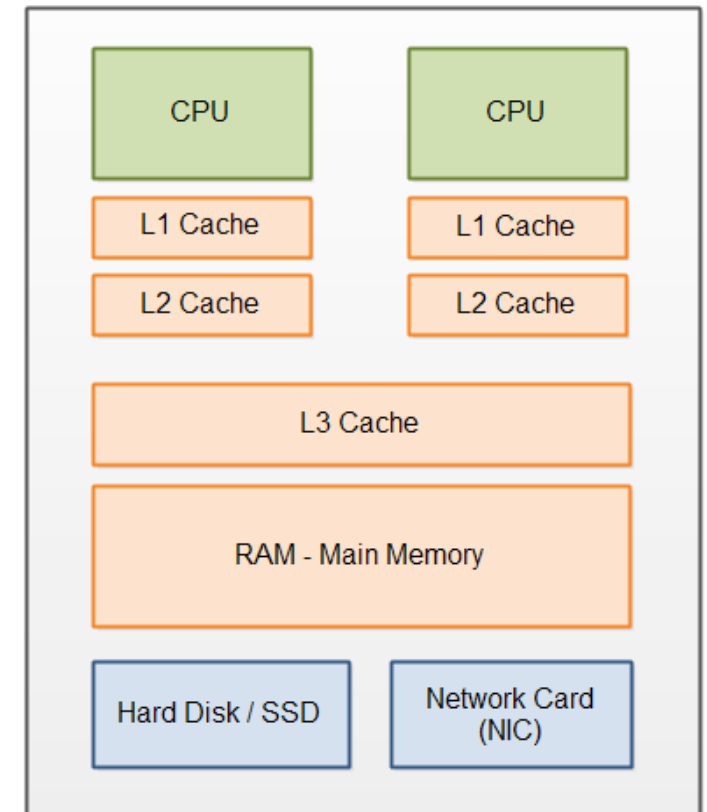
Exploiting Locality

- If the system accessed some data at a given time, it is likely that it will access the same or 'nearby' data 'soon'
 - Nearby => spatial, e.g., watching a new episode on iPlayer in the UK
 - Soon => temporal, e.g., next instruction in a program
- Improve access time constraints, reduce load on paths to main server.
- Caching: store copies of frequently accessed data from a **larger, slower storage** in **small, fast storage instances**.
 - Example: Filesystem caching - get the speed of RAM and the capacity of disk



Hierarchy

- Recursive decomposition of a system into smaller pieces that depend only on the parent for proper execution
- Examples: Computer memory, Addressing in Networks
- No single point of control
- Highly scalable
- Leaf-to-leaf communication can be expensive
 - shortcuts help



Binding and indirection

- **Abstraction** is good
 - Allows generality of description: Domain names, variable pointers
- **Binding**: translate an abstraction to an instance
- **Indirection**: use a reference to **access something**, instead of directly
 - Provides **flexibility** because the reference can change without affecting the higher-level system
 - Dynamic binding allows easy optimisation for current conditions
 - E.g., use domain-name to get the closest instance of a service
- Common technique in the Internet; can direct user to the best server based on current conditions

A records

IPv4 address

✓ 142.250.185.238



Google LLC

Location [Frankfurt am Main, Hesse, Germany](#)

AS AS15169

AS name Google LLC

<https://www.nslookup.io/domains/google.com/dns-records/#thenetherlands>

A records

IPv4 address

✓ 142.250.71.78



Google LLC

Location [Sydney, New South Wales, Australia](#)

AS AS15169

AS name Google LLC



<https://www.nslookup.io/domains/google.com/dns-records/#australia>

Virtualization

- A combination of indirection and multiplexing
- Create a **virtual version of a resource** (hardware, OS, network, or storage) so that multiple independent systems or applications can share it as if each has its own dedicated resource
- Examples:
 - virtual memory
 - virtual modem
 - Cloud computing
 - Santa claus
- Can cleanly and dynamically reconfigure the system
- Simplifies programming for developers; you can code without worrying synchronization with other programs

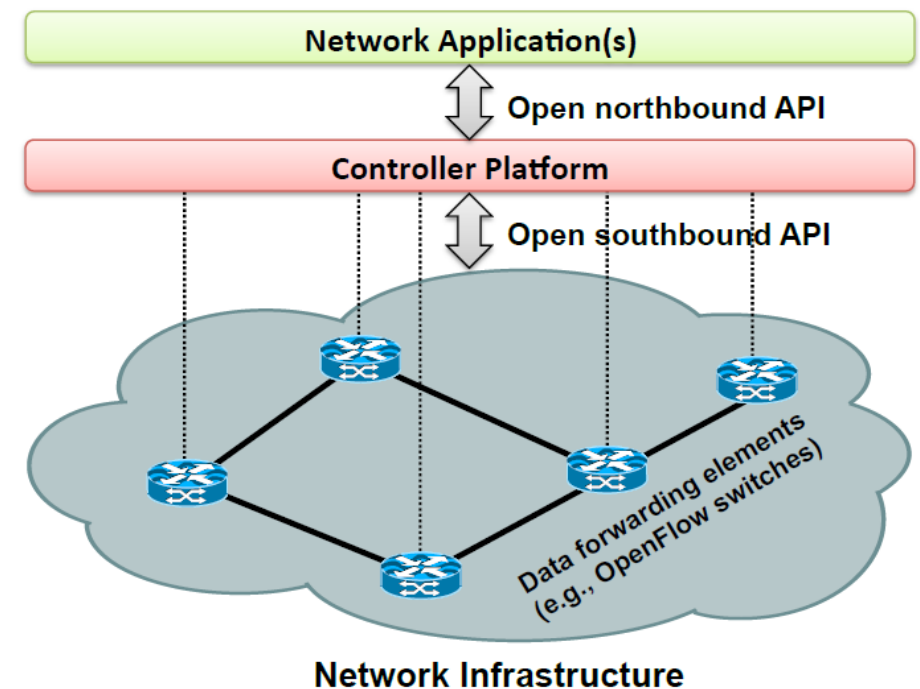
Randomization

- Allows us to break a tie fairly
- Can also be used as a simple mechanism to spread load
- Examples
 - resolving contention in a broadcast medium
 - Use hashes on load balancers to spread traffic across backend servers.
- Statistics are important.
 - If task effort is not the same, fairness can be compromised



Separating data and control

- Divide actions that happen once per data transfer (control plane) from actions that occur once per packet (data plane)
 - **Data plane** = “*what is being carried*” (payload, instructions, user content).
 - **Control plane** = “*how and where it should go*” (rules, scheduling, routing).
- Can increase throughput by optimizing processing in the data path
- Can improve intelligence through the data plane
- Example
 - Software-defined networks
- On the other hand, keeping control information in data plane can improve optimisation
 - per-packet QoS



Performance analysis and tuning

- Use the techniques discussed to tune existing systems
- Steps
 - measure
 - characterize workload
 - build a system model
 - analyse
 - implement



Conclusions

- System design is the science of combining multiple components to deliver specific functionalities.
 - Optimize a metrics, by using resources and meeting constraints.
- We use design patterns to improve the performance of a system.
 - Explore system design patterns in action as we discuss in more depth networking and OS technologies.
- Next Topic: Operating Systems and Processes

Labs start next week, check your calendar!!!!

Questions (1)

- From your experience with everyday life, can you find examples that use system design techniques to improve system performance?
 - E.g. cashiers in the supermarket multiplex queues, TAs in the lab allow support indirection

Question (2)

- Assume that Database1 and Database2 serve the same static content to different parts of the network. Can you think of an application of a design pattern in the architecture below to improve the system cost? What impact does your design have on the system security?

